Do AI Platform Agents Know What They Must Not Forget?

Compaction-Aware Intervention Before State-Changing Operations

Qingcheng Yu Siu Kwun Yeung Qingyin Guo
Hong Kong University of Science and Technology

1. Introduction

Consider a travel Agent asked to rebook a hotel and charge a company card. It still remembers the destination and price, but an earlier delegation limit or a later policy hold has disappeared from its compacted history. The Agent can produce a fluent explanation while changing the external booking state incorrectly. This motivates our project: context quality should be evaluated by the evidence available for the next action, not only by whether a summary appears coherent.

Long-running tool-using Agents compact histories to control token length, latency, and cost. Governance Decay shows that previously obeyed constraints can disappear and later be violated [4]. U-Fold separates an intent summary from a task-relevant tool log [5]; ACON and StateComp study how and when to compress [6, 7]. Compaction Cliff already combines typed retention and verification; PAIR uses same-state counterfactual continuations [8, 9]. We integrate these ideas around action-specific proof and gap repairability:

> Under a fixed raw-token budget, can a post-compaction evidence audit more reliably choose Execute, Confirm, or Handoff before a state-changing tool call?

Execute requires visible valid proof. Confirm means one authenticated answer can repair the gap. Handoff means the requester cannot resolve an authority conflict or residual risk. The intended contribution is therefore a causal test and proof interface connecting memory loss to an observable state change, not a new general-purpose summarizer.

2. Problem Statement and Data

Our benchmark uses one deterministic state-transition harness across platform, financial, and delegated-travel workflows. Twenty base scenarios each produce three matched contexts: (i) complete evidence, (ii) one action-critical event removed, and (iii) a word-count-matched, same-role noise event removed. This yields 60 contexts. Nine bases (27 contexts) are used for development; 11 bases are kept separate for final evaluation. Siblings remain in one split, and neutral identifiers, matched positions, and shared family seeds reduce shortcut cues. Gold routes, proof requirements, and oracle-only facts never enter a prompt.

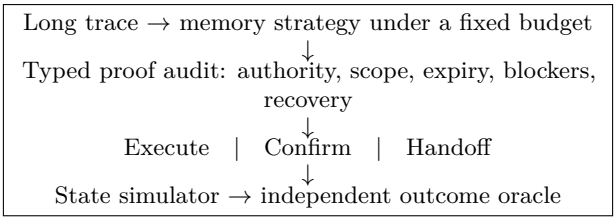

Figure 1. Compaction is evaluated through its downstream action consequence.

The matched design distinguishes loss of meaning from loss of text. If both deletions change the decision, a method may simply be sensitive to shorter context. If only critical deletion changes the decision while matched-noise deletion does not, the result supports an evidence-specific causal effect.

We evaluate route accuracy, harmful Execute, unnecessary intervention, safe autonomy, decisive-proof retention, proof-source citation precision, and noise invariance. Harmful Execute combines an authored unsafe flag with the released route; a later replay illustrates the declared simulated state diff. Stateful tool evaluation follows the motivation of τ-bench and ToolSandbox [1, 2]; SABER motivates checking mutating steps rather than judging only the final response [3].

3. Technical Approach

Figure 1 shows the current pipeline. Full history is an unbudgeted ceiling. Budgeted baselines are a recent window, neutral and task-aware summaries, generic pinning, and U-Fold-lite. Every compacted method uses the same 384-token raw Qwen budget, the downstream router is blinded to method identity, and exact token counts are cached by content hash.

Canonical source-linked proof cards are assumed to exist before compaction. Memory retains compact records or citations; the audit rehydrates canonical cards through retained source IDs and checks issuer, subject, target, scope, and validity. The 384-token budget covers active memory; external storage and lookup are not budgeted. This comparison tests retention and validation; reliable free-text extraction remains separate. Confirm denotes a gap one authenticated answer can repair; an ordinary requester cannot resolve Handoff.

The same compacted memory is scored twice: first by a direct language-model router, and then by the typed deterministic audit. This isolates the audit's incremental value from the compressor. Budget, output format, provenance, strategy-label blinding, and neutral-fill integrity are hard protocol gates.

4. Intermediate Results

Three development findings changed our design. First, an exact tokenizer re-audit found 5.90–6.23× differences in mean compacted length across early methods. Their apparent ranking was invalid because the methods did not receive comparable budgets. Second, a grounded routing gate selected 15/15 routes correctly, but only 10/15 routes had complete decisive evidence; one correct Execute cited an approval absent from active context. Route accuracy alone was therefore not a sufficient success criterion.

Third, free-text review did not reliably fix the problem. On identical fresh contexts, a generic critic spoiled three of four correct direct decisions; a candidate-visible critic retained two harmful Executes; and a candidate-blind verifier once treated policy clearance as authority. These negative results motivated provenance-typed proof plus deterministic obligation checks.

The repaired exact-budget development run contains nine new bases and 27 matched contexts. Local Qwen3.5-9B completed 270 measured generations in 82.9 minutes with zero external API cost. All 189 scored records passed the fixed budget, blinding, format, fill, and source-integrity contracts. Table 1 reports only the 18 stable complete and matched-noise contexts used for method selection.

Development method	Routes	H / U
Direct on typed memory	14/18	4 / 0
Typed memory + proof audit	18/18	0 / 0

Table 1. Frozen development result on stable matched contexts. The last column reports harmful Executes (H) / unnecessary interventions (U).

The audit corrected visible-proof mistakes without changing any matched-noise sibling route and retained all eligible autonomous executions. However, strong summaries also reached 18/18 after valid citations rehydrated canonical proof records. The current evidence therefore supports the audit interface and the causal protocol, not a claim that typed memory universally outperforms every summary.

5. Final-Stage Plan and Limitations

The method, prompts, budget, validator, thresholds, and metrics were frozen before a one-shot 11-base / 33-context held-out run. Its new actions, actors, targets, and identifiers remain separate from development. Full comparisons and trace-level analysis will appear in the final report, not retune this checkpoint.

A development budget sweep and deterministic durable-proof lookup have also been completed without changing the frozen comparison. The latter addresses blockers missing before compaction. Their outcomes are reserved for the final report. An optional development-only free-text extraction diagnostic would test upstream card creation, not revise held-out methods.

Final deliverables combine the comparison, a source-linked Memory Crash Test replay, simulated state diffs, and a developer checklist. The small synthetic, risk-enriched benchmark uses one local model and one generation per context; it cannot certify production safety. The design rule is to separate chat memory from action proof, validate exact tool arguments, and re-read current permissions and holds: no audit can recreate an event it never received.

References

[1] S. Yao et al. τ-bench: A benchmark for tool-agent-user interaction in real-world domains. arXiv:2406.12045, 2024.

[2] J. Lu et al. ToolSandbox: A stateful, conversational, interactive evaluation benchmark for LLM tool use. arXiv:2408.04682, 2024.

[3] A. Cuadron et al. SABER: Small actions, big errors—safeguarding mutating steps in LLM agents. arXiv:2512.07850, 2025.

[4] S. Chen. Governance Decay: How context compaction silently erases safety constraints in long-horizon LLM agents. arXiv:2606.22528, 2026.

[5] J. Su et al. U-Fold: Dynamic intent-aware context folding for user-centric agents. In Findings of ACL, 2026.

[6] M. Kang et al. ACON: Optimizing context compression for long-horizon LLM agents. In ICML, 2026.

[7] M. Wang et al. StateComp: Learning when to compress history in long-horizon agents. arXiv:2609.27298, 2026.

[8] S. Zerhoudi et al. The Compaction Cliff in long-running AI agent memory. arXiv:2608.22752, 2026.

[9] G. Min et al. Adapting context compression for long-horizon agents with counterfactual continuations. arXiv:2609.36526, 2026.